

GAMM: Generalized Additive Mixed Models

GAM.jl Contributors

Table of contents

Introduction	1
Model specification	2
Setup	2
Example 1: Gaussian GAMM with Random Intercepts	2
Data	2
Fitting with <code>gamm()</code>	3
Random effects	3
Variance components	4
Comparison with true values	4
Visualizing the Gaussian GAMM	5
Equivalence with <code>s(subject, bs=:re)</code>	5
Example 2: Poisson GAMM for Count Data	6
Fitting	6
Random effects	7
Visualizing the Poisson GAMM	7
Example 3: Alternative Formula Interfaces	8
Using <code>@formula</code> with <code>(1 group)</code>	8
Using <code>@formula</code> with <code>re(group)</code>	9
Consistency	9
Prediction	9
Summary	10

Introduction

A **Generalized Additive Mixed Model** (GAMM) combines the smooth, nonlinear covariate effects of a GAM with **grouped random effects** — random intercepts and slopes that vary across levels of a grouping factor (e.g., subjects, sites, or time series). GAMMs are essential when:

- Observations are **clustered** (repeated measures on subjects, spatial replicates at sites)
- Between-cluster variability needs to be **explicitly modeled** (not just absorbed by smooth terms)
- You want to separate **population-level smooth trends** from **cluster-specific deviations**

In penalized regression terms, a random intercept $b_j \sim N(0, \sigma_b^2)$ is equivalent to adding a ridge penalty on the group indicator columns — which is precisely what `s(group, bs=re)` already does in a standard GAM. GAM.jl's `gamm()` provides the familiar mixed-model syntax (`1|group`) as a convenient front-end to this same machinery.

Model specification

For a response y_{ij} (observation i in group j):

$$g(\mu_{ij}) = \mathbf{x}_i^\top \beta + \sum_k f_k(x_{ik}) + b_j, \quad b_j \sim N(0, \sigma_b^2)$$

where g is the link function, f_k are smooth functions estimated via penalized regression splines, and b_j are random intercepts.

Setup

```
using GAM
using CSV
using DataFrames
using Distributions
using GLM: LogLink, IdentityLink
using Statistics: mean, std, var, cor
using StatsAPI: fitted, deviance, nobs, coef, predict
using Plots
using Printf
```

Example 1: Gaussian GAMM with Random Intercepts

Data

Simulated data with 12 subjects, each observed 40 times. The population-level signal is $\mu(x) = 1.5 \sin(1.5x)$, with subject-specific random intercepts ($\sigma_b = 0.6$) and residual noise ($\sigma_\varepsilon = 0.4$).

```
dat = CSV.read("data_gaussian_gamm.csv", DataFrame)
println("n = $(nrow(dat)), subjects = $(length(unique(dat.subject)))")
println("y range: [$(round(minimum(dat.y); digits=2)), $(round(maximum(dat.y); digits=2))]")
```

```
n = 480, subjects = 12
y range: [-2.41, 2.38]
```

Fitting with gamm()

The (1 | subject) syntax specifies a random intercept for each subject level:

```
m = gamm(@gamm_formula(y ~ s(x, k=15) + (1 | subject)), dat)
println(m)
```

Generalized Additive Mixed Model

Family: Normal

Link: IdentityLink

Fixed Effects Coefficients:

```
[1] = 0.025752
```

Smooth Terms:

```
s(x,bs=tp) edf = 7.82
```

Random Effects:

```
(1 | subject) = 0.2206 (n_levels = 12)
```

Deviance: 78.5108

REML: 282.9764

Scale est.: 0.169240

n = 480

Random effects

Extract per-subject random intercept estimates (BLUPs) using `ranef()`:

```

re = ranef(m)
est = vec(re.subject.effects)
levels = re.subject.levels
for (lev, eff) in zip(levels, est)
  @printf(" Subject %2s: b = %+.3f\n", lev, eff)
end

```

```

Subject 1: b = -0.090
Subject 2: b = -0.033
Subject 3: b = +0.163
Subject 4: b = -0.233
Subject 5: b = -0.080
Subject 6: b = -0.086
Subject 7: b = +0.152
Subject 8: b = +0.281
Subject 9: b = +0.060
Subject 10: b = -0.129
Subject 11: b = +0.249
Subject 12: b = -0.257

```

Variance components

VarCorr() returns the estimated random effect standard deviation:

```

vc = VarCorr(m)
for v in vc
  @printf(" %s: = %.4f (n_levels = %d)\n", v.label, v.std, v.n_levels)
end
residual_scale = m isa GAM.GammModel ? m.gam_model.scale : m.scale
@printf(" Residual: = %.4f\n", sqrt(residual_scale))

```

```

(1 | subject): = 0.2206 (n_levels = 12)
Residual: = 0.4114

```

Comparison with true values

```

true_re = [dat.re_true[findfirst(dat.subject .== s)] for s in sort(unique(dat.subject))]
@printf("Correlation of estimated vs true RE: %.4f\n", cor(est, true_re))

```

Correlation of estimated vs true RE: 0.8471


Visualizing the Gaussian GAMM

```
# Population smooth on a prediction grid (unknown subject → zero RE)
x_grid = range(minimum(dat.x), maximum(dat.x); length=200)
pop_pred = predict(m, DataFrame(x=x_grid, subject=fill(999, 200)))

p1 = scatter(dat.x, dat.y; group=dat.subject, markersize=2, alpha=0.5,
             xlabel="x", ylabel="y", title="Gaussian GAMM: data by subject",
             legend=:none)
plot!(p1, x_grid, pop_pred; color=:black, linewidth=3, label="population smooth")

n_groups = length(levels)
p2 = bar(1:n_groups, [est true_re]; label=["Estimated" "True"], legend=:topright,
        xlabel="Subject", ylabel="Random intercept",
        title="Random intercepts: estimated vs true",
        xticks=(1:n_groups, string.(levels)))

plot(p1, p2; layout=(1, 2), size=(900, 400))
```



Equivalence with `s(subject, bs=:re)`

In GAM.jl, `gamm(@gamm_formula(y ~ s(x) + (1|subject)), ...)` is mathematically equivalent to `gam(@gam_formula(y ~ s(x) + s(subject, bs=:re)), ...)`. Both treat the random intercept as a smooth with identity penalty:

```

m_gam = gam(@gam_formula(y ~ s(x, k=15) + s(subject, bs=:re)), dat)
@printf("Fitted values correlation: %.6f\n", cor(fitted(m), fitted(m_gam)))
@printf("Scale (gamm): %.6f\n", m.gam_model.scale)
@printf("Scale (gam): %.6f\n", m_gam.scale)

```

```

Fitted values correlation: 0.999995
Scale (gamm): 0.169240
Scale (gam): 0.169392

```

Example 2: Poisson GAMM for Count Data

Count data with 8 sites, each observed 60 times. The true log-rate has a smooth trend plus site-specific random intercepts ($\sigma_b = 0.4$):

$$\log(\lambda_{ij}) = 1 + 0.8 \sin(x_i) + b_j, \quad b_j \sim N(0, 0.16)$$

```

dat2 = CSV.read("data_poisson_gamm.csv", DataFrame)
println("n = $(nrow(dat2)), sites = $(length(unique(dat2.site)))")
println("y range: [$(minimum(dat2.y)), $(maximum(dat2.y))]" )

```

```

n = 480, sites = 8
y range: [0.0, 15.0]

```

Fitting

Pass `Poisson()` as the family — the canonical `LogLink` is used automatically:

```

m2 = gamm(@gamm_formula(y ~ s(x, k=15) + (1 | site)), dat2, Poisson())
println(m2)

```

Generalized Additive Mixed Model

```

Family: Poisson
Link:    LogLink

```

```

Fixed Effects Coefficients:
[1] = 1.139613

```

```

Smooth Terms:
  s(x,bs=tp)          edf =    5.31

Random Effects:
  (1 | site)          =    0.1150  (n_levels = 8)

Deviance:          553.9613
REML:              942.2983
Scale est.:        1.000000
n = 480

```

Random effects

```

re2 = ranef(m2)
est2 = vec(re2.site.effects)
true_re2 = [dat2.re_true[findfirst(dat2.site .== s)] for s in sort(unique(dat2.site))]
@printf("RE correlation with truth: %.4f\n", cor(est2, true_re2))

vc2 = VarCorr(m2)
@printf("Estimated _RE: %.4f (true: 0.4)\n", vc2[1].std)

```

```

RE correlation with truth: 0.6420
Estimated _RE: 0.1150 (true: 0.4)

```

Visualizing the Poisson GAMM

```

# Population smooth on prediction grid
x_grid2 = range(minimum(dat2.x), maximum(dat2.x); length=200)
pop_pred2 = predict(m2, DataFrame(x=x_grid2, site=fill(999, 200)))

site_levels = sort(unique(dat2.site))
p1 = scatter(dat2.x, dat2.y; group=dat2.site, markersize=2, alpha=0.5,
  xlabel="x", ylabel="y (count)", title="Poisson GAMM: data by site",
  legend=:none)
plot!(p1, x_grid2, pop_pred2; color=:black, linewidth=3, label="population mean")

n_sites = length(site_levels)
p2 = bar(1:n_sites, [est2 true_re2]; label=["Estimated" "True"], legend=:topright,
  xlabel="Site", ylabel="Random intercept (log-scale)",


```

```

    title="Site random effects: estimated vs true",
    xticks=(1:n_sites, string.(site_levels)))

plot(p1, p2; layout=(1, 2), size=(900, 400))

```



Example 3: Alternative Formula Interfaces

GAM.jl supports multiple ways to specify random effects:

Using @formula with (1|group)

StatsModels.jl parses $(1|\text{group})$ as a `FunctionTerm{typeof(|)}`, which `gamm()` detects automatically:

```

_scale(m) = m isa GAM.GammModel ? m.gam_model.scale : m.scale
m3a = gamm(@formula(y ~ cr(x, 15) + (1|subject)), dat)
@printf("@formula path: scale = %.6f\n", _scale(m3a))

```

```
@formula path: scale = 0.169240
```


Using @formula with re(group)

The `re()` convenience function provides an alternative syntax:

```
m3b = gamm(@formula(y ~ cr(x, 15) + re(subject)), dat)
@printf("re() path: scale = %.6f\n", _scale(m3b))
```

```
re() path: scale = 0.175579
```

Consistency

All three formula interfaces produce the same fit:

```
m3c = gamm(@gamm_formula(y ~ s(x, k=15) + (1|subject)), dat)
@printf("cor(@gamm_formula, @formula(|)): %.6f\n", cor(fitted(m3c), fitted(m3a)))
@printf("cor(@gamm_formula, @formula(re)): %.6f\n", cor(fitted(m3c), fitted(m3b)))
```

```
cor(@gamm_formula, @formula(|)): 1.000000
cor(@gamm_formula, @formula(re)): 0.997444
```

Prediction

Predict on new data — unknown group levels get zero random effect contribution:

```
# Known subjects
df_known = DataFrame(x=[0.0, 0.5, 1.0], subject=[1, 2, 3])
pred_known = predict(m, df_known)
@printf("Predictions (known subjects): [%.3f, %.3f, %.3f]\n", pred_known...)

# New subject (never seen)
df_new = DataFrame(x=[0.0, 0.5, 1.0], subject=[999, 999, 999])
pred_new = predict(m, df_new)
@printf("Predictions (new subject): [%.3f, %.3f, %.3f]\n", pred_new...)
```

```
Predictions (known subjects): [0.393, 0.801, 0.495]
Predictions (new subject): [0.495, 0.650, 0.544]
```

Summary

Feature	Syntax
Random intercept	<code>(1 \ group)</code> or <code>re(group)</code>
Gaussian family	<code>gamm(formula, data)</code>
Non-Gaussian	<code>gamm(formula, data, Poisson())</code>
Equivalent GAM	<code>gam(@gam_formula(y ~ s(x) + s(group, bs=:re)), ...)</code>
Random effects	<code>ranef(m)</code>
Variance components	<code>VarCorr(m)</code>
Prediction	<code>predict(m, newdata)</code>

GAM.jl's `gamm()` delegates to the same proven PIRLS+REML machinery used by `gam()`, treating random effects as smooth terms with identity penalty matrices. This means GAMM results are numerically equivalent to the corresponding `s(group, bs=:re)` formulation, but with the convenience of mixed-model notation.